

From voice calls to resolved cases.

A simple strategy for building Voni in regulated industries.

The short answer: ElevenLabs can handle the conversation. Voni should make sure the work is allowed, recorded, and finished.

CROWDED MARKET

Voice agent builder

Prompts, voices, calls, workflows, tools, and analytics.



RECOMMENDED CATEGORY

Governed case resolution

Policy, permission, evidence, deadlines, people, and outcomes.

“ElevenLabs runs the conversation. Voni governs the case the conversation is allowed to advance.”

Do not compete with ElevenLabs on generic features.

ElevenLabs already offers a strong horizontal platform: voices, languages, phone calls, WhatsApp, workflows, tools, knowledge bases, transfers, testing, analytics, guardrails, versioning, privacy controls, data residency, and private deployment.

LET PROVIDERS HANDLE

The live conversation

- ✓ Speech and voices
- ✓ Turn-taking and interruption
- ✓ Session workflows and tools
- ✓ Phone transport and transfers
- ✓ Transcripts and call analytics

VONI MUST HANDLE

The institution's obligation

- ✓ What case is open
- ✓ Which policy version applies
- ✓ What facts are verified
- ✓ Who may approve an action
- ✓ Whether the case is truly resolved

Provider

Conversation

"Your document has expired."



Voni

Case

Missing proof of address, due Friday.



Voni

Safe action

Send a secure link after confirmation.



Source system

Outcome

Replacement received and reconciled.

The test for differentiation

If removing Voni leaves the institution with the same case state, permission controls, evidence, and resolution process, Voni has not built a defensible product.

Keep the safety work. Change what the product is for.

KEEP AND STRENGTHEN

Safety foundations

Organization boundaries, durable background jobs, calling rules, consent checks, confirmed voice actions, and the signed-in copilot.

TRANSFORM

Current objects

Agents become runtime assets. Campaigns open cases. Leads become people and verified contacts. Calls and messages become evidence.

STOP LEADING WITH

Generic breadth

A blank agent builder, a general workflow canvas, a voice catalogue, dialer breadth, generic RAG, and call-volume charts.

Seven simple product objects

1. Policy pack

The approved rule book: exact wording, allowed actions, calling rules, sources, dates, jurisdiction, reviewers, and signed versions.

2. Case

One record across calls, messages, documents, systems, deadlines, missing items, and responsible people.

3. Permission gate

A clear answer to: is this person or system allowed to do this action now?

4. Evidence log

The policy clause, facts, statement, approval, tool result, and time behind every important step.

5. Resolution plan

A fixed checklist for what must happen, with freedom in how the conversation explains it.

6. Specialist queue

A place for disputes, uncertainty, vulnerability, approvals, failures, and overdue work.

7. Provider adapter

A standard connection that lets Voni change voice, phone, or model suppliers without losing the case.

Start with proactive application resolution.

This means contacting a person when an application is stuck because one clear, fixable item is missing or out of date.

Application resolution

83

Start here. Clear trigger, reversible actions, measurable delay, and a clean human decision boundary.

Difficult moments

57

Useful later. Detecting vulnerability is uncertain, so the first role should be careful escalation to a specialist.

Collections

55

Not first. The value is real, but conduct risk, disputes, and irreversible harm make it a poor opening proof.

Why the first workflow wins

Narrow

One known missing item

Measurable

Time to resolution

Reversible

Explain, send, schedule

Safe

No approval decision

Reusable

Banking, insurance, government

Voni may

- ✓ Explain the exact missing item
- ✓ Send a secure upload link after confirmation
- ✓ Schedule a callback
- ✓ Update a contact preference after readback

A human must

- ✗ Waive a requirement
- ✗ Accept or reject a document
- ✗ Approve, price, or score an application
- ✗ Make fraud, AML, sanctions, or adverse decisions

Aisha's application is stuck.

Her proof of address has expired. She prefers Arabic and cannot visit a branch this week.

ILLUSTRATIVE BANKING CASE

09:00

1 Voni opens a case

The bank system reports one missing requirement. Voni pins the approved policy and starts the deadline.

10:14

2 The conversation verifies only what is needed

Aisha proves enough identity to discuss the missing item. The agent still cannot discuss or decide the full application.

10:17

3 Aisha confirms one safe action

The agent explains the effect in Arabic, reads it back, and asks permission to send a secure link to her verified number.

10:26


4 Voni checks the real outcome

The bank system confirms receipt. Voni records the proof and closes the missing item, not the application.

Human boundary: if Aisha disputes the date, the records conflict, or the document needs judgment, the case goes to a specialist.

The value is not “a call was completed.” The value is that the requirement was corrected faster, the action was permitted, and an auditor can reconstruct what happened.

Every important action follows the same path.



Provider independence is a design choice.

01	Institution controls	Organizations, roles, retention, residency, rollout, rollback, and evaluation.
02	Policy and permission	Signed rules, identity level, risk, expiry, approvers, and dual control.
03	Case resolution	Cases, evidence, plans, deadlines, queues, and reconciliation.
04	Provider adapters	Common events for speech, phone, language, transfer, transcript, and quality.
05	Replaceable suppliers	ElevenLabs, AssemblyAI, carriers, models, identity services, and source systems.

Slow work must survive the page. Policy checks, data sync, exports, and provider changes use Voni's durable job system. Live calls stay in the foreground.

Voni configures. The institution approves.

Start with a managed implementation. Do not offer unrestricted self-service policy editing until the same workflow has been repeated safely.

Institution policy owner	Owns the source policy, intended meaning, approved wording, effective dates, and final sign-off.
Operations owner	Owns the current process, service levels, queues, staffing, integrations, and baseline results.
Compliance, legal, privacy, security	Review jurisdiction, data, access, retention, provider, and launch conditions.
Voni team	Structures the policy pack, configures the workflow, connects providers, writes tests, maps evidence, and monitors the pilot.
Technology team or partner	Connects source systems, networks, identity services, and local operations.

Voni may prepare a policy version. Only an authorized institution reviewer can approve it. The institution always owns the customer outcome.

The same pattern works elsewhere.

INSURANCE

Omar's claim lacks one photograph

Voni can explain the missing angle, send a secure request, and track receipt. A claims employee decides coverage, fraud concerns, and settlement.

GOVERNMENT

Maya's permit lacks a tenancy record

Voni can explain the published requirement and track the deadline. The authority decides eligibility and accepts exceptions.

Earn the platform one phase at a time.

- | | | |
|---|--|--|
| 0 | 4-6 weeks
Prove the problem | Map one real policy and historical case set. Establish baseline outcomes, authority, queues, integrations, and stop rules. |
| 1 | 8-12 weeks
Run a controlled pilot | Build signed policy versions, cases, one safe action, one source connection, one provider connection, a specialist queue, and full evidence. |
| 2 | 3-6 months
Harden and repeat | Add a second provider, rollback, retention, dual control, a second missing item, and a second language only after quality matches. |
| 3 | 6-12 months
Reuse the pattern | Expand into insurance, permits, licenses, benefits, and claims. Certify partners only after repeated proof. |

Measure

- ✓ Corrected without a visit
- ✓ Time to resolution
- ✓ Repeat contact
- ✓ Human handling time
- ✓ Policy mistakes and blocked unsafe actions
- ✓ Quality by language
- ✓ Time to reconstruct an audit

Stop if

- ✗ Buyers only want a generic voice agent
- ✗ No partner can supply approved policy and integration access
- ✗ Institutions will not pay for control and evidence
- ✗ Source systems cannot confirm outcomes
- ✗ Complaints, exceptions, or unequal language outcomes increase

Charge for governed work, not minutes.

Commercial hypothesis

- ✓ Annual control-plane fee
- ✓ Active or resolved case volume
- ✓ Implementation and policy-pack setup
- ✓ Optional language, evaluation, and audit packages
- ✓ Voice usage passed through or contracted directly

A lighthouse partner must provide

- ✓ An approved policy
- ✓ Anonymized historical cases
- ✓ Reliable integration access
- ✓ Specialists for escalations
- ✓ Baseline performance data
- ✓ A measurable expansion or stop decision

Questions to answer before building

1. Which missing item causes the most correctable delay?
2. Can the source system send a trusted trigger and final receipt?
3. What may be revealed at each identity level?
4. Is the call service, marketing, or another category?
5. What consent, disclosure, recording, and contact rules apply?
6. Which wording must be exact in each language?
7. What makes a case disputed, vulnerable, or unsafe?
8. Who may approve each action?
9. Which system wins when records conflict?
10. What evidence do auditors and complaints teams need?
11. What are today's resolution, effort, complaint, and language results?
12. Where will specialists do their work?
13. Which provider and residency choices are acceptable?
14. What result means expand, pause, or stop?

Recommended first sale: a Voni-managed pilot for one missing application requirement, with the institution approving policy and retaining every consequential decision.

Primary sources

The comparison uses current ElevenLabs documentation. Regulatory sources show representative patterns, not universal compliance.

ElevenLabs: ElevenAgents overview

ElevenLabs: agent workflows

ElevenLabs: structured procedures

ElevenLabs: MCP tools and approvals

ElevenLabs: agent testing

ElevenLabs: versioning

ElevenLabs: guardrails

ElevenLabs: analytics

ElevenLabs: SIP trunking

ElevenLabs: batch calls

ElevenLabs: privacy

ElevenLabs: data residency

ElevenLabs: private deployment

CBUAE: Consumer Protection Standards

CBUAE: telemarketing controls

UAE: Cabinet Resolution No. 56 of 2024

European Commission: AI Act framework

European Union: GDPR

UK FCA: Consumer Duty

Canada: Directive on Automated Decision-Making

US FCC: AI-generated voices in robocalls

W3C: WCAG 2.2

NIST: AI Risk Management Framework

ISO/IEC 42001

FINAL RECOMMENDATION

Build governed case resolution.

Start with one application problem. Let ElevenLabs remain a voice supplier. Make Voni the place where policy, permission, evidence, people, and outcomes come together.